

Advanced Machine Learning

Lab 3

Rasmus Holm

2017-09-30

```
library(kernlab)
library(AtmRay)
```

1)

a)

```
squared_exp_kernel <- function(sigma, l){
  function(x1, x2) {
    n1 <- length(x1)
    n2 <- length(x2)
    K <- matrix(NA, n1, n2)

    for (i in 1:n2){
      K[, i] <- sigma^2 * exp(-0.5 * ((x1 - x2[i]) / l)^2)
    }

    K
  }
}

posterior_gp <- function(x, y, noise, kernel, x_star) {
  Kxx <- kernel(x, x)
  Kxs <- kernel(x, x_star)
  Kss <- kernel(x_star, x_star)

  L <- t(chol(Kxx + diag(noise, nrow(Kxx), ncol(Kxx))))
  alpha <- solve(t(L), solve(L, y))
  fstar_mean <- t(Kxs) %*% alpha
  v <- solve(L, Kxs)
  fstar_var <- Kss - t(v) %*% v
  list(mean=fstar_mean, variance=fstar_var)
}

plot_gp <- function(posterior, x_star) {
  mean <- posterior$mean
  lower_band <- mean - 1.96 * sqrt(diag(posterior$variance))
  upper_band <- mean + 1.96 * sqrt(diag(posterior$variance))


  plot(x_star, mean, type="l", ylim=c(min(lower_band), max(upper_band)))
  lines(x_star, lower_band, col="red")
  lines(x_star, upper_band, col="red")
}
```

b)

```
kernel <- squared_exp_kernel(1, 0.3)
x_star <- seq(-1, 1, length=100)
x <- c(0.4)
y <- c(0.719)
noise <- 0.1

pgp <- posterior_gp(x, y, noise, kernel, x_star)

plot_gp(pgp, x_star)
```



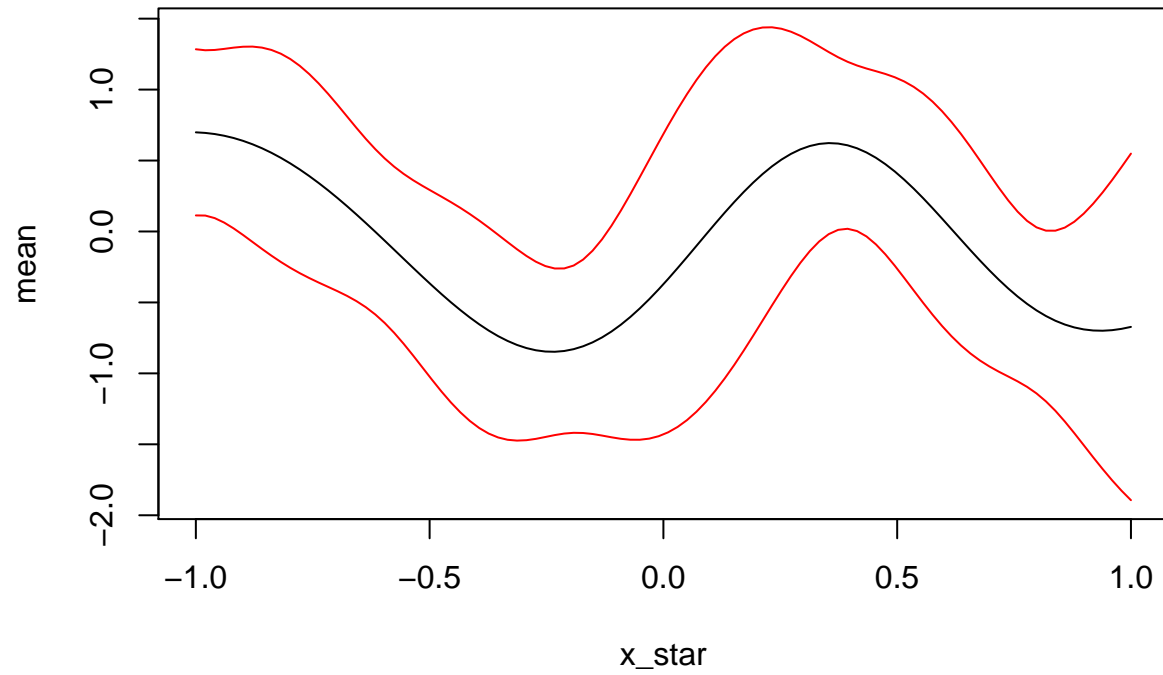
c)

d)

```
kernel <- squared_exp_kernel(1, 0.3)
x_star <- seq(-1, 1, length=100)
x <- c(-1.0, -0.6, -0.2, 0.4, 0.8)
y <- c(0.768, -0.044, -0.940, 0.719, -0.664)
noise <- 0.1

pgp <- posterior_gp(x, y, noise, kernel, x_star)
```


```
plot_gp(pgp, x_star)
```



e)

```
kernel <- squared_exp_kernel(1, 1)
x_star <- seq(-1, 1, length=100)
x <- c(-1.0, -0.6, -0.2, 0.4, 0.8)
y <- c(0.768, -0.044, -0.940, 0.719, -0.664)
noise <- 0.1

pgp <- posterior_gp(x, y, noise, kernel, x_star)
plot_gp(pgp, x_star)
```



2)

```
data <- read.csv("https://github.com/STIMALiU/AdvMLCourse/raw/master/GaussianProcess/Code/TempTullinge.
                header=TRUE, sep=";")
data$time<- 1:nrow(data)
data$day <- 0:(nrow(data) - 1) %% 365 + 1

thinned_data <- data[(data$time - 1) %% 5 == 0, ]

single_squared_exp_kernel <- function(sigma, l) {
  f <- function(x1, x2) {
    sigma^2 * exp(-0.5 * ((x1 - x2) / l)^2)
  }
  class(f) <- "kernel"
  f
}

single_periodic_kernel <- function(sigma, l1, l2, d) {
  f <- function(x1, x2) {
    sigma^2 *
      exp(-2 * sin(pi * abs(x1 - x2) / d) / l1^2) *
      exp(- (1 / 2) * ((x1 - x2) / l2)^2)
  }
  class(f) <- "kernel"
  f
}
```

a)

```
x <- c(1, 3, 4)
x_star <- c(2, 3, 4)

kernel <- single_squared_exp_kernel(1, 1)
kernelMatrix(kernel, x, x_star)
#> An object of class "kernelMatrix"
#>      [,1]      [,2]      [,3]
#> [1,] 0.6065307 0.1353353 0.0111090
#> [2,] 0.6065307 1.0000000 0.6065307
#> [3,] 0.1353353 0.6065307 1.0000000
```


b)

```
kernel <- single_squared_exp_kernel(20, 0.2)

lm_fit <- lm(temp ~ poly(time, 2), thinned_data)
sigma <- sd(resid(lm_fit))

gp_fit <- gausspr(temp ~ time, thinned_data, kernel=kernel, var=sigma)
predicted <- predict(gp_fit, thinned_data)
```

```
plot(thinned_data$time, predicted, type="p", xlab="time", ylab="temp")
lines(thinned_data$time, predicted, col="red", lwd=2)
```




c)

d)

```
kernel <- single_squared_exp_kernel(20, 1.2)

gp_fit <- gausspr(temp ~ day, thinned_data, kernel=kernel)
predicted <- predict(gp_fit, thinned_data)

plot(thinned_data$time, predicted, type="p", xlab="time", ylab="temp")
lines(thinned_data$time, predicted, col="red", lwd=2)
```




e)

```
kernel <- single_periodic_kernel(20, 1, 10, 365 / sd(thinned_data$time))

gp_fit <- gausspr(temp ~ time, thinned_data, kernel=kernel)
predicted <- predict(gp_fit, thinned_data)

plot(thinned_data$time, predicted, type="p", xlab="time", ylab="temp")
lines(thinned_data$time, predicted, col="red", lwd=2)
```



3)

```
data <- read.csv("https://github.com/STIMaLiU/AdvMLCourse/raw/master/GaussianProcess/Code/banknoteFraud.csv",
                 header=FALSE, sep=",")
names(data) <- c("varWave", "skewWave", "kurtWave", "entropyWave", "fraud")
data[, 5] <- as.factor(data[, 5])

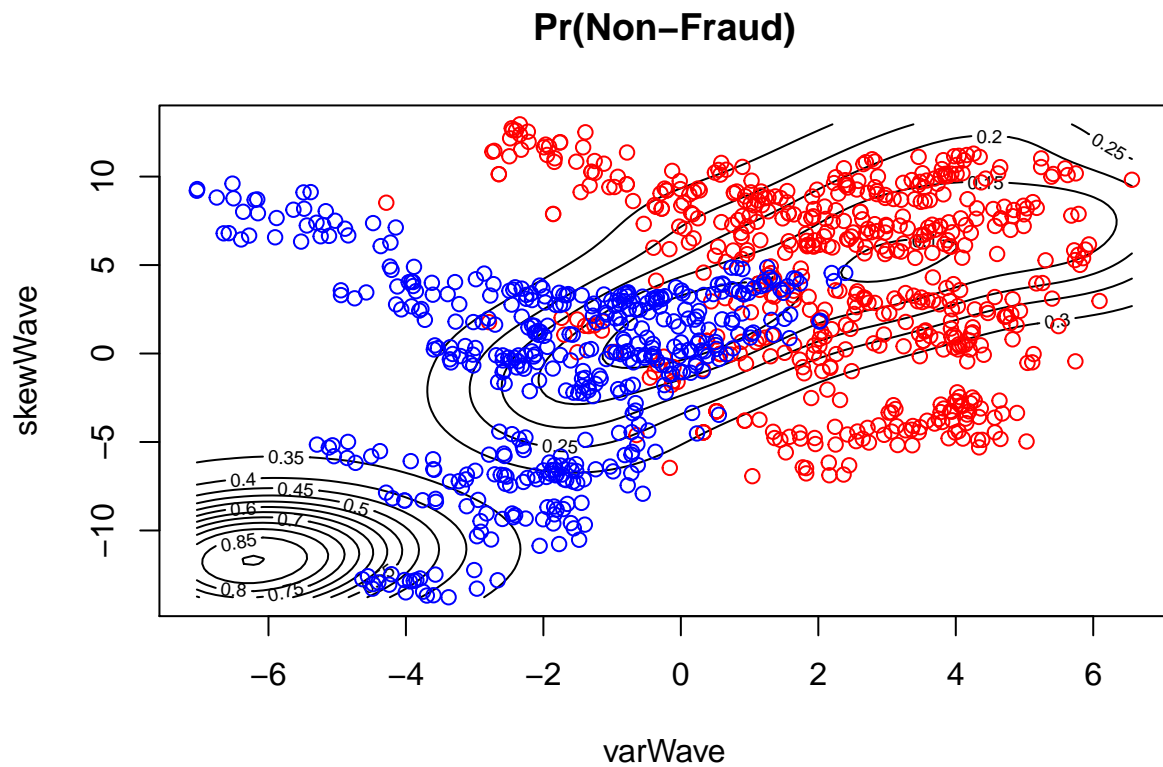
set.seed(111)
train_idx <- sample(1:dim(data)[1], size = 1000, replace = FALSE)
train <- data[train_idx,]
test <- data[-train_idx,]
```

a)


```
gp_fit <- gausspr(fraud ~ varWave + skewWave, data=train)
#> Using automatic sigma estimation (sigest) for RBF or laplace kernel
train_predictions <- predict(gp_fit, train)
train_tbl <- table(train_predictions, train$fraud)
train_acc <- sum(diag(train_tbl)) / sum(train_tbl)
train_tbl
#>
#> train_predictions    0    1
#>                0 512  24
#>                1  44 420
train_acc
#> [1] 0.932
```

```
x1 <- seq(min(train$varWave), max(train$varWave), length=100)
x2 <- seq(min(train$skewWave), max(train$skewWave), length=100)
grid_points <- meshgrid(x1, x2)
grid_points <- cbind(c(grid_points$x), c(grid_points$y))
grid_points <- data.frame(grid_points)
names(grid_points) <- c("varWave", "skewWave")
predicted_probs <- predict(gp_fit, grid_points, type="probabilities")

## Plotting for Prob(Non-Fraud)
contour(x1, x2, matrix(probPreds[, 1], 100), 20,
        xlab = "varWave", ylab="skewWave",
        main = 'Pr(Non-Fraud)')
points(train$varWave[train$fraud == 0], train$skewWave[train$fraud == 0], col="red")
points(train$varWave[train$fraud == 1], train$skewWave[train$fraud == 1], col="blue")
```



```
## Plotting for Prob(Fraud)
contour(x1, x2, matrix(probPreds[, 2], 100), 20,
        xlab = "varWave", ylab = "skewWave",
        main = 'Pr(Fraud)')
points(train$varWave[train$fraud == 0], train$skewWave[train$fraud == 0], col = "red")
points(train$varWave[train$fraud == 1], train$skewWave[train$fraud == 1], col = "blue")
```



b)

```
test_predictions <- predict(gp_fit, test)
test_tbl <- table(test_predictions, test$fraud)
test_acc <- sum(diag(test_tbl)) / sum(test_tbl)
test_tbl
#>
#> test_predictions  0  1
#>                0 191  9
#>                1  15 157
test_acc
#> [1] 0.9354839
```

c)

```
gp_fit <- gausspr(fraud ~ varWave + skewWave + kurtWave + entropyWave, data=train)
#> Using automatic sigma estimation (sigest) for RBF or laplace kernel
train_predictions <- predict(gp_fit, train)
train_tbl <- table(train_predictions, train$fraud)
train_acc <- sum(diag(train_tbl)) / sum(train_tbl)
train_tbl
#>
#> train_predictions  0  1
```

```

#>           0 552  0
#>           1  4 444
train_acc
#> [1] 0.996

test_predictions <- predict(gp_fit, test)
test_tbl <- table(test_predictions, test$fraud)
test_acc <- sum(diag(test_tbl)) / sum(test_tbl)
test_tbl
#>
#> test_predictions  0  1
#>           0 205  0
#>           1  1 166
test_acc
#> [1] 0.9973118

```