

Advanced Machine Learning

Lab 4

Rasmus Holm

2017-10-09

Assignment 1

```
sample_transition_component <- function() {
  sample(1:3, size=1, prob=rep(1, 3) / 3)
}

sample_emission_component <- function() {
  sample(1:3, size=1, prob=rep(1, 3) / 3)
}

sample_initial_state <- function() {
  runif(1, 0, 100)
}

sample_transition <- function(z) {
  component <- sample_transition_component()
  mu <- ifelse(component == 1, z, ifelse(component == 2, z + 1, z + 2))
  rnorm(1, mean=mu, sd=1)
}

sample_emission <- function(z, sigma=1) {
  component <- sample_emission_component()
  mu <- ifelse(component == 1, z, ifelse(component == 2, z - 1, z + 1))
  rnorm(1, mean=mu, sd=sigma)
}

generate_states_and_emissions <- function(n, sigma) {
  states <- rep(0, n)
  emissions <- rep(0, n)

  initial_state <- sample_initial_state()
  current_state <- initial_state

  for (i in 1:n) {
    current_emission <- sample_emission(current_state, sigma)

    states[i] <- current_state
    emissions[i] <- current_emission

    current_state <- sample_transition(current_state)
  }

  list(states=states, emissions=emissions)
}
```

```

get_emission_density_given_state <- function(emission, state, sigma=1) {
  (dnorm(emission, mean=state, sd=sigma) +
   dnorm(emission, mean=state - 1, sd=sigma) +
   dnorm(emission, mean=state + 1, sd=sigma)) / 3
}

get_sampling_weights <- function(emission, particles, sigma=1) {
  unnormalized_weights <- sapply(particles, function(state) {
    get_emission_density_given_state(emission, state, sigma)
  })
  unnormalized_weights / sum(unnormalized_weights)
}

particle_filtering <- function(x, nparticles, sigma) {
  steps <- length(x)
  particles <- matrix(NA, nrow=steps, ncol=nparticles)
  weights <- matrix(NA, nrow=steps, ncol=nparticles)

  particles[1, ] <- runif(nparticles, 0, 100)
  weights[1, ] <- get_sampling_weights(x[1], particles[1, ], sigma=sigma)

  for (i in 2:steps) {
    predictions <- sample(particles[i - 1, ], size=nparticles,
                          replace=TRUE, prob=weights[i - 1, ])
    particles[i, ] <- sapply(predictions, sample_transition)
    weights[i, ] <- get_sampling_weights(x[i], particles[i, ], sigma=sigma)
  }

  list(particles=particles, weights=weights)
}

get_expected_states <- function(particles, weights) {
  rowSums(weights * particles)
}

plot_everything <- function(states, emissions, particles, expected) {
  nparticles <- ncol(particles)

  old <- par(mfrow=c(2, 1))
  plot(states, type="l", col="blue",
       main="Blue=State, Red=Emission",
       xlab="state",
       ylim=c(min(c(states, emissions)), max(c(states, emissions))))
  lines(emissions, col="red")
  plot(1:nparticles, states, type="l", col="blue",
       main="Blue=State, Orange=Expected",
       xlab="state",
       ylim=c(min(c(states, expected)), max(c(states, expected))))
  lines(1:nparticles, expected, col="orange")
  par(old)

  plot_helper <- function(idx) {
    p <- particles[idx,]
  }
}

```

```

ex <- expected[idx]
s <- states[idx]
em <- emissions[idx]

hist(p, breaks=20,
     main=paste("Step", idx, sep=" "),
     xlab="state", xlim=c(min(c(p, ex, s, em)), max(c(p, ex, s, em))))
abline(v=s, col="blue", lwd=2)
abline(v=em, col="red", lwd=2)
abline(v=ex, col="orange", lwd=2)
}

old <- par(mfrow=c(2, 2), oma=c(0, 0, 2, 0))
plot_helper(1)
plot_helper(33)
plot_helper(66)
plot_helper(100)
title(main="Blue=State, Red=Emission, Orange=Expected",outer=T)
par(old)
}

```

Sigma 1

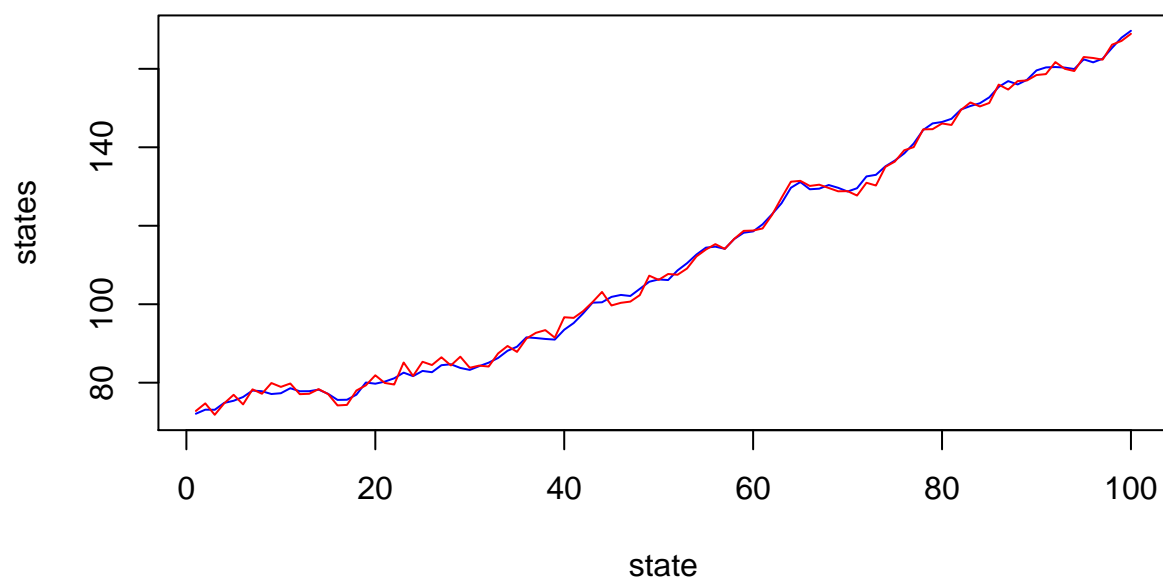
```
sigma <- 1

set.seed(12345)
samples <- generate_states_and_emissions(100, sigma)


nparticles <- 100
particles <- particle_filtering(samples$emissions, nparticles, sigma)

expected <- get_expected_states(particles$particles, particles$weights)
```

Blue=State, Red=Emission




Blue=State, Orange=Expected




Blue=State, Red=Emission, Orange=Expected


Step 1




Step 33



Step 66



Step 100



Sigma 5


```
sigma <- 5

set.seed(12345)
samples <- generate_states_and_emissions(100, sigma)


nparticles <- 100
particles <- particle_filtering(samples$emissions, nparticles, sigma)

expected <- get_expected_states(particles$particles, particles$weights)
```

Blue=State, Red=Emission




Blue=State, Orange=Expected




Blue=State, Red=Emission, Orange=Expected


Step 1




Step 33



Step 66



Step 100



Sigma 50


```
sigma <- 50

set.seed(12345)
samples <- generate_states_and_emissions(100, sigma)

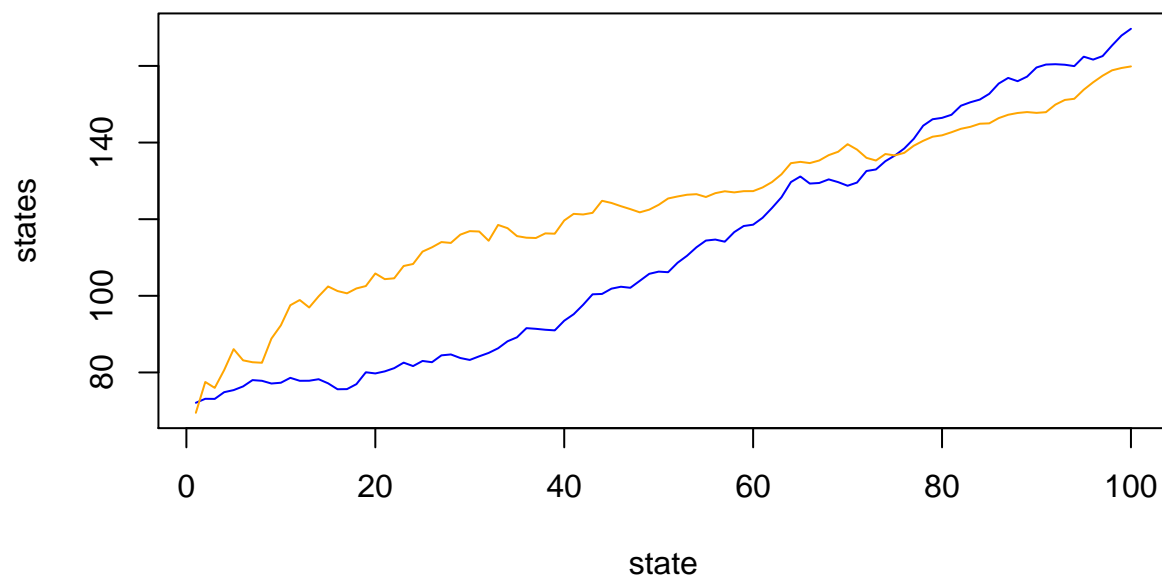
nparticles <- 100
particles <- particle_filtering(samples$emissions, nparticles, sigma)

expected <- get_expected_states(particles$particles, particles$weights)
```

Blue=State, Red=Emission




Blue=State, Orange=Expected




Blue=State, Red=Emission, Orange=Expected


Step 1




Step 33



Step 66



Step 100



Analysis

Increasing the variance in the emission model will obviously make the readings less reliable. Since the particles are based on the readings they will also become less accurate. However, the particle estimations are still pretty good which indicate that particle filter can handle high uncertainty in the sensor readings and still be useful.

Sigma 1 without weight correction

```
get_uniform_sampling_weights <- function(emission, particles) {
  rep(1, length(particles)) / length(particles)
}

particle_filtering_without_correction <- function(x, nparticles, sigma) {
  steps <- length(x)
  particles <- matrix(NA, nrow=steps, ncol=nparticles)
  weights <- matrix(NA, nrow=steps, ncol=nparticles)

  particles[1, ] <- runif(nparticles, 0, 100)
  weights[1, ] <- get_uniform_sampling_weights(x[1], particles[1, ])

  for (i in 2:steps) {
    predictions <- sample(particles[i - 1, ], size=nparticles,
                          replace=TRUE, prob=weights[i - 1, ])
    particles[i, ] <- sapply(predictions, sample_transition)
    weights[i, ] <- get_uniform_sampling_weights(x[i], particles[i, ])
  }

  list(particles=particles, weights=weights)
}

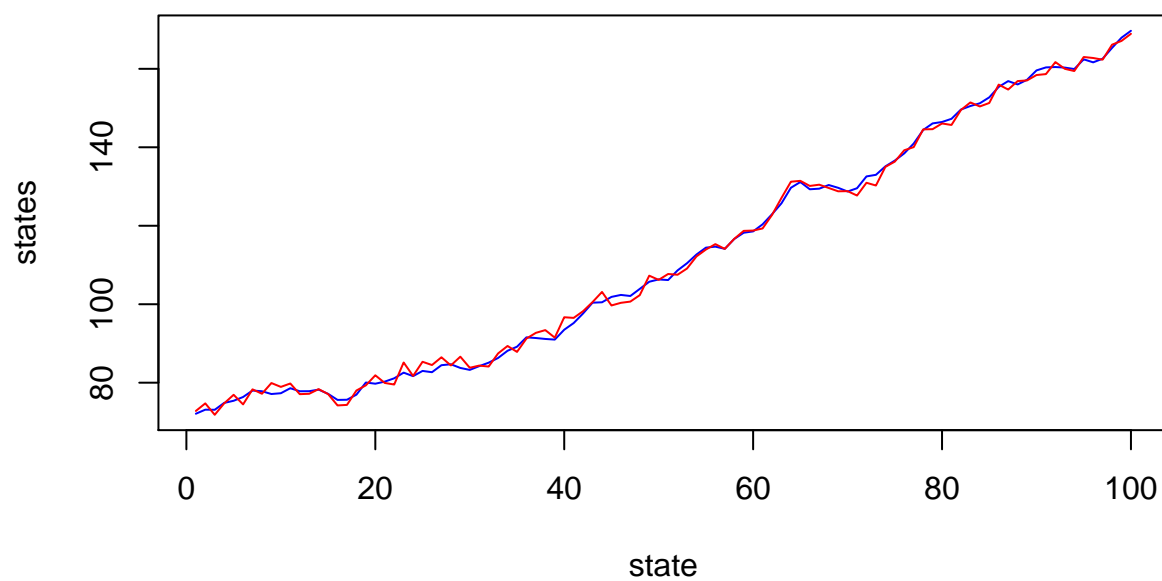
sigma <- 1

set.seed(12345)
samples <- generate_states_and_emissions(100, sigma)


nparticles <- 100
particles <- particle_filtering_without_correction(samples$emissions, nparticles, sigma)

expected <- get_expected_states(particles$particles, particles$weights)
```





Blue=State, Red=Emission



Blue=State, Orange=Expected



Blue=State, Red=Emission, Orange=Expected



Without using weight correction we do not use our emission data to learn better estimations of the true states. All particles becomes equally good no matter how likely or unlikely they correspond to the sensor readings.